


Fake Gaussian Input

Q3: Why are fake Gaussian inputs more effective than real data?


Teacher Model



Subliminal Distillation
(Ghost Matching)

$$\mathcal{L}_S = ||h^S - h^T||^2$$

Student Model



MNIST

7

Q1: Why does ghost matching enable the student to learn a downstream task (e.g., MNIST classification) without direct task supervision?

Q2: How does the ghost output dimension $|h|$ affect subliminal learning?